

## Exercise # 2. Iterative Methods For Linear Systems.

Alexandre Rodrigues (2039952)

January 8, 2022

### Question 1

Using as a test the example usage, with  $tol = 1 \times 10^{-8}$  and limiting the iterations to  $maxit = 250$  I got the following results.




Figure 1: Residual norm vs iteration number for PCG methods,  $maxit = 250$

Method	Iterations	Final Residual	Computational Time
Matlab PCG without preconditioning	183	$1.9591 \times 10^{-7}$	0.077s
Matlab PCG IC(0)	78	$1.5293 \times 10^{-7}$	0.068s
My PCG implementation	250	$2.3 \times 10^{-3}$	0.151s

Table 1: Results of PCG methods,  $maxit = 250$

When  $maxit$  is large enough to guarantee convergence in all implementations we get the following results:




Figure 2: Residual norm vs iteration number for PCG methods,  $maxit = 750$

Method	Iterations	Final Residual	Computational Time
Matlab PCG without preconditioning	183	$1.9591 \times 10^{-7}$	0.054s
Matlab PCG IC(0)	78	$1.5293 \times 10^{-7}$	0.063s
My PCG implementation	720	$9.6833 \times 10^{-9}$	0.351s

Table 2: Results of PCG methods,  $maxit = 750$

My implementation is slower to converge but produces better final residual values.

## Question 2

The spectral condition number of A is

$$\kappa(A) = \frac{\lambda_{max}(A)}{\lambda_{min}(A)}. \quad (1)$$

In Matlab I used the `condest(A)` function to estimate the condition number of a sparse matrix A.

$n_x$	$h$	$\kappa(A)$	$\sqrt{\kappa(A)}$	CG	PCG(0)	PCG( $10^{-2}$ )	PCG( $10^{-3}$ )
102	$1.0000 \times 10^{-4}$	$6.0107 \times 10^3$	77.5288	283	87	45	17
202	$2.5000 \times 10^{-5}$	$2.3810 \times 10^4$	154.3039	532	159	78	30
402	$6.2500 \times 10^{-6}$	$9.4770 \times 10^4$	307.8473	948	282	137	53
802	$1.5625 \times 10^{-6}$	$3.7814 \times 10^5$	614.9304	1792	533	258	97

Table 3: Iterations of PCG methods for each value of  $n_x$  and respective values of  $h$  and  $\kappa(A)$

One can note from the table the dependence of the number of iterations on  $h = \frac{1}{N} = \frac{1}{(nx-2)^2}$ . The number of iterations is halved when  $n_x$  approximately doubles.

## Question 3

show theoretically ??

When using the Cholesky preconditioner with no fill in, I didn't get the expected results. Both Matlab's and my implementation converged in only one iteration.




Figure 3: Residual norm vs iteration number for PCG methods with  $IC(0)$  preconditioner

Due to the bad results, I tried to remove preconditioning from my implementation by setting  $L$  as the identity matrix,  $L = \text{speye}(\text{size}(L))$ .

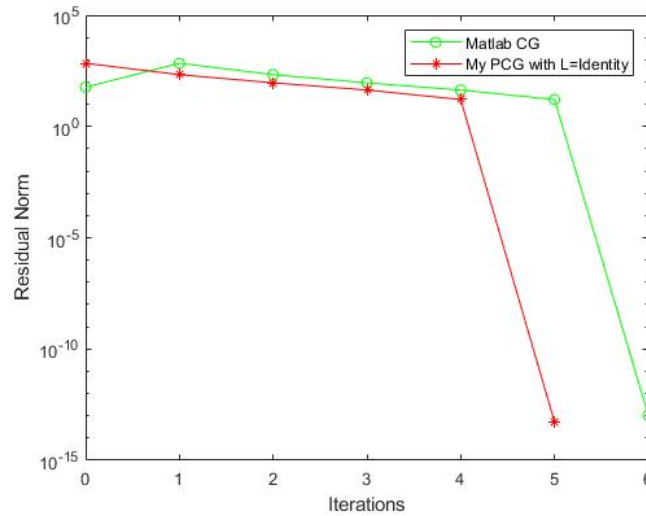


Figure 4: Residual norm vs iteration number for PCG methods without preconditioning

Method	Iterations	Final Residual	Computational Time
Matlab PCG	6	$9.2128 \times 10^{-14}$	0.021s
My PCG	5	$1.2744 \times 10^{-13}$	0.012s

Table 4: Results for each value of implementation, no preconditioning

These results show the theoretical calculations, my implementation is still better than expected.

## Question 4




Figure 5: Residual norm vs iteration number for PCG methods without preconditioning

Preconditioner	Iterations	Final Residual	Computational Time
None	253	$1.1367 \times 10^{-4}$	0.254s
Jacobi	36	$9.3198 \times 10^{-5}$	0.053s
IC(0)	10	$1.1155 \times 10^{-4}$	0.046s

Table 5: Results for each preconditioner

There is a very clear improvement when using preconditioning. It is also noticeable the superior characteristics of the incomplete Cholesky preconditioner relative to Jacobi.

## Question 5




Figure 6: Residual norm vs iteration number for GMRES methods

Method	Iterations	Final Residual	Computational Time
Matlab GMRES	527	$1.2073 \times 10^{-12}$	9.097s
My GMRES	509	$1.2231 \times 10^{-12}$	10.032s

Table 6: Results for each GMRES implementation

These results show that the methods have very similar convergence characteristics. My implementation has a smaller number of iterations but the other results are slightly worse than the ones achieved with Matlab's implementation.

## Question 6




Figure 7: Residual norm vs iteration number for preconditioned GMRES methods

Method	Iterations	Final Residual	True Residual	Computational Time
Matlab GMRES	38	$1.5592 \times 10^{-10}$	$4.5350 \times 10^{-13}$	0.121s
My GMRES	39	$2.3797 \times 10^{-13}$	$7.1893 \times 10^{-14}$	4.943s

Table 7: Results for each preconditioned GMRES implementation

There are clear differences in the residuals and computation time values. My implementation is 40 times slower but produces a true residual 5 times smaller.

Solving the linear system in 3b produced the same unexpected results as in that question. When using the Cholelsky precontionier with no fill in both Matlab's and my implementation converged in only one iteration.

Method	Iterations	Final Residual	Computational Time
GMRES	1	$3.8481 \times 10^{-16}$	0.027s
My PCG	1	$1.7554 \times 10^{-17}$	0.003s

Table 8: Iterations for each value of nx

When I removed preconditioning form my implementation by setting  $L$  as the identity matrix,  $L = \text{speye}(\text{size}(L))$ .

Method	Iterations	Final Residual	Computational Time
GMRES	6	$3.0413 \times 10^{-14}$	0.102s
My PCG	5	$1.0468 \times 10^{-13}$	0.012s

Table 9: Results for each value of implementation, no preconditioning




Figure 8: Residual norm vs iteration number for GMRES methods without preconditioning

As in 3b, these results show the theoretical calculations, my implementation is still better than expected.

## Question 7

restart	Iterations	Final Residual	Computational Time
10	1149	$1.9901 \times 10^{-12}$	1.735s
20	739	$1.9741 \times 10^{-12}$	1.443s
30	88	$1.3800 \times 10^{-12}$	0.242s
50	41	$9.9416 \times 10^{-13}$	0.135s

Table 10: Results for each value of `restart`




Figure 9: Residual norm vs iteration number for each value of `restart`

One can notice a clear improvement in convergence with the increase of the `restart` value.

...

## Question 8

Using `maxit=550`, `tol=1e-8`...

Tolerance	Iterations	Prec. Time	Tsol	Ttotal	Final Residual	$\rho$
$2 \times 10^{-2}$	1316	37.92s	44.53s	82.45s	$5.2065 \times 10^{-7}$	0.4537
$1 \times 10^{-2}$	4444	40.55s	22.86s	63.42s	$5.7213 \times 10^{-7}$	0.5807
$3 \times 10^{-3}$	150	51.78s	7.39s	59.17s	$6.4998 \times 10^{-7}$	0.9401
$1 \times 10^{-3}$	67	47.74s	3.63s	51.37s	$8.5337 \times 10^{-7}$	1.4544
$1 \times 10^{-4}$	26	43.30s	2.30s	45.60s	$9.1517 \times 10^{-7}$	3.5140
$1 \times 10^{-5}$	12	96.73s	2.69s	99.42s	$9.4359 \times 10^{-7}$	9.0720

Table 11: Results for each value of tolerance




Figure 10: Residual norm vs iteration number for each value of tolerance